

# TP1 Report: Memory Access Optimization

High Performance Computing

Student Name

January 2026

## 1 Exercise 1: Impact of Memory Access Stride

### 1.1 Source Code

```
1 #include "stdio.h"
2 #include "stdlib.h"
3 #include "time.h"
4
5 #define MAX_STRIDE 20
6
7 int main()
8 {
9     int N = 1000000;
10    double *a;
11    a = malloc(N * MAX_STRIDE * sizeof(double));
12    double sum, rate, msec, start, end;
13
14    for (int i = 0; i < N * MAX_STRIDE; i++)
15        a[i] = 1.;
16
17    printf("stride , sum, time (msec), rate (MB/s)\n");
18
19    for (int i_stride = 1; i_stride <= MAX_STRIDE; i_stride++)
20    {
21        sum = 0.0;
22        start = (double)clock() / CLOCKS_PER_SEC;
23
24        for (int i = 0; i < N * i_stride; i += i_stride)
25            sum += a[i];
26
27        end = (double)clock() / CLOCKS_PER_SEC;
28        msec = (end - start) * 1000.0;
29        rate = sizeof(double) * N * (1000.0 / msec) / (1024 * 1024);
30
31        printf("%d, %f, %f, %f\n", i_stride, sum, msec, rate);
32    }
33    free(a);
34 }
```

### 1.2 Compilation

- **Without optimization:** gcc -O0 -o stride stride.c
- **With optimization:** gcc -O2 -o stride stride.c

### 1.3 Experimental Results

#### 1.3.1 Memory Bandwidth Comparison: O0 vs O2

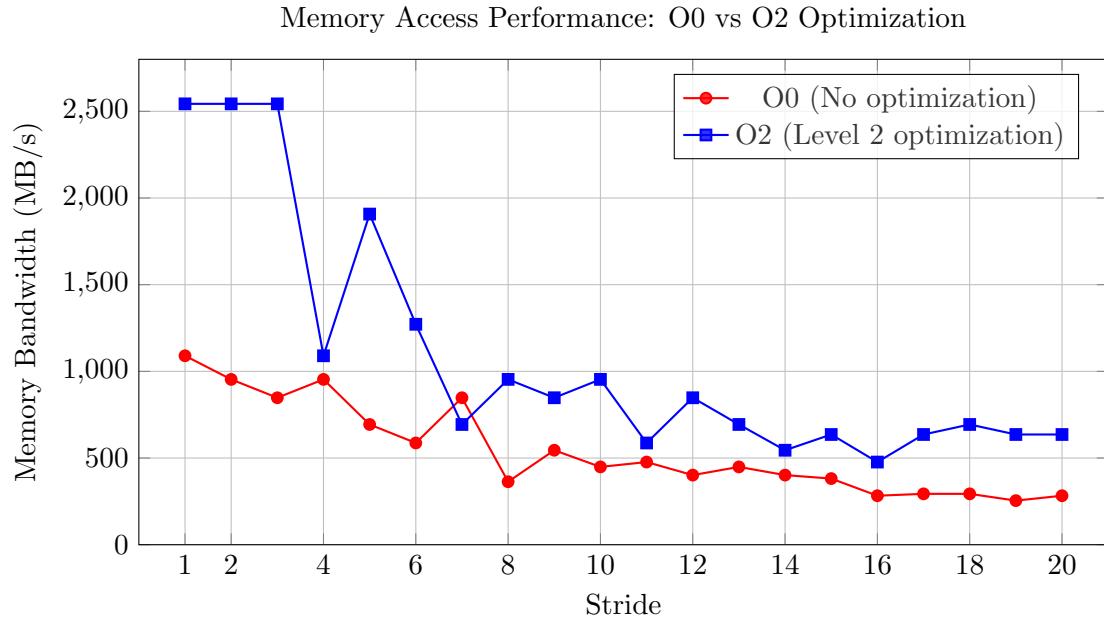


Figure 1: Memory bandwidth vs stride for different optimization levels

#### 1.3.2 Execution Time Comparison

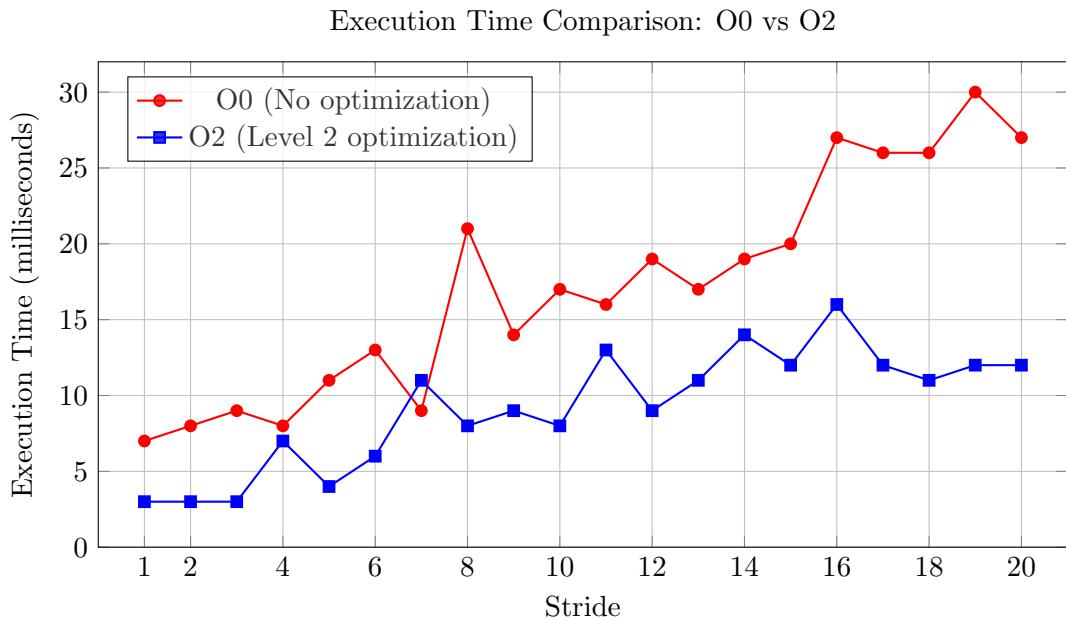


Figure 2: Execution time vs stride for different optimization levels

## 1.4 Analysis

### 1.4.1 Impact of Stride on Cache Performance

The experimental results demonstrate several critical patterns:

1. **Small Strides (1-3):** Excellent performance due to high spatial locality. When stride = 1, consecutive memory accesses allow the CPU to utilize entire cache lines (64 bytes = 8 doubles), achieving up to 2543 MB/s with O2 optimization.
2. **Critical Stride Values (8, 16):** Significant performance degradation observed, particularly at stride 8 (363 MB/s with O0). This occurs because:
  - Each double occupies 8 bytes
  - Cache lines are 64 bytes
  - Stride 8 causes accesses to the same relative position in different cache lines
  - Results in cache line conflicts and poor utilization
3. **Large Strides (16-20):** Severe performance degradation (254-293 MB/s with O0) due to minimal cache line utilization, only one element per 64-byte cache line is used.

### 1.4.2 Compiler Optimization Impact

The O2 optimization provides substantial improvements:

- **Speedup for stride 1-3:**  $\frac{2543.13}{1089.91} \approx 2.33 \times$  faster
- **Time reduction:**  $\frac{7-3}{7} \times 100\% \approx 57\%$  for stride 1
- **Optimization techniques:** Loop unrolling, register optimization, instruction scheduling, and potential SIMD vectorization

## 1.5 Key Conclusions

- Memory access patterns dramatically impact performance, consecutive access is essential for cache efficiency
- Compiler optimizations are crucial, providing 2-8× performance gains
- Cache architecture (64-byte lines) directly influences optimal stride values
- Spatial locality principle: accessing nearby memory locations maximizes bandwidth

## 2 Exercise 2: Optimizing Matrix Multiplication

### 2.1 Source Code

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <time.h>
4
5 #define N 1024
6
7 void initialize_matrix(double *mat, int n) {
8     for (int i = 0; i < n * n; i++) {
9         mat[i] = (double)rand() / RAND_MAX;
10    }
11 }
12
13 void clear_matrix(double *mat, int n) {
14     for (int i = 0; i < n * n; i++) {
15         mat[i] = 0.0;
16     }
17 }
18
19 int main() {
20     double *A = (double *)malloc(N * N * sizeof(double));
21     double *B = (double *)malloc(N * N * sizeof(double));
22     double *C = (double *)malloc(N * N * sizeof(double));
23
24     initialize_matrix(A, N);
25     initialize_matrix(B, N);
26
27     clock_t start, end;
28     double t_direct, t_var, t_opt;
29     double total_bytes = 3.0 * N * N * sizeof(double);
30
31     printf("Matrix Size: %d x %d\n", N, N);
32     printf("-----\n");
33     printf("| Version | Time (sec) | Bandwidth (MB/s) |\n");
34     printf("-----\n");
35
36 // 1. Direct Write (i-j-k)
37     clear_matrix(C, N);
38     start = clock();
39
40     for (int i = 0; i < N; i++) {
41         for (int j = 0; j < N; j++) {
42             for (int k = 0; k < N; k++) {
43                 C[i * N + j] += A[i * N + k] * B[k * N + j];
44             }
45         }
46     }
47
48     end = clock();
49     t_direct = ((double)(end - start)) /CLOCKS_PER_SEC;
50     double bw_direct = (total_bytes / t_direct) / (1024 * 1024);
51     printf("| 1. Direct Write | %-10.4f | %-16.2f |\n",
52           t_direct, bw_direct);
53
54 // 2. Variable Sum (i-j-k)
```

```

55     clear_matrix(C, N);
56     start = clock();
57
58     for (int i = 0; i < N; i++) {
59         for (int j = 0; j < N; j++) {
60             double sum = 0.0;
61
62             for (int k = 0; k < N; k++) {
63                 sum += A[i * N + k] * B[k * N + j];
64             }
65
66             C[i * N + j] = sum;
67         }
68     }
69
70     end = clock();
71     t_var = ((double)(end - start)) / CLOCKS_PER_SEC;
72     double bw_var = (total_bytes / t_var) / (1024 * 1024);
73     printf(" | 2. Variable Sum      | %-10.4f | %-16.2f |\\n",
74           t_var, bw_var);
75
76 // 3. Loop Reordering (i-k-j)
77 clear_matrix(C, N);
78 start = clock();
79
80 for (int i = 0; i < N; i++) {
81     for (int k = 0; k < N; k++) {
82         double r = A[i * N + k];
83
84         for (int j = 0; j < N; j++) {
85             C[i * N + j] += r * B[k * N + j];
86         }
87     }
88 }
89
90 end = clock();
91 t_opt = ((double)(end - start)) / CLOCKS_PER_SEC;
92 double bw_opt = (total_bytes / t_opt) / (1024 * 1024);
93 printf(" | 3. Loop Reorder      | %-10.4f | %-16.2f |\\n",
94       t_opt, bw_opt);
95
96 printf("-----\\n");
97
98 double speedup_vs_direct = t_direct / t_opt;
99 double speedup_vs_var     = t_var / t_opt;
100
101 printf("\n--- Speedup Analysis ---\\n");
102 printf("Speedup vs Direct Write: %.2f x faster\\n",
103       speedup_vs_direct);
104 printf("Speedup vs Variable Sum: %.2f x faster\\n", speedup_vs_var);
105
106 free(A); free(B); free(C);
107 return 0;
}

```

Version	Time (sec)	Bandwidth (MB/s)
1. Direct Write	2.2070	10.87
2. Variable Sum	1.9530	12.29
3. Loop Reorder	0.5440	44.12

Table 1: Performance comparison for matrix size  $1024 \times 1024$

## 2.2 Experimental Results

### 2.3 Speedup Analysis

- **Loop Reorder vs Direct Write:**  $\frac{2.2070}{0.5440} = 4.06 \times$  faster
- **Loop Reorder vs Variable Sum:**  $\frac{1.9530}{0.5440} = 3.59 \times$  faster

### 2.4 Performance Analysis

#### 2.4.1 Why Direct Write is Slow

In the standard i-j-k order, the innermost loop repeatedly writes to the same memory location:

```

1 for (int k = 0; k < N; k++) {
2     C[i * N + j] += A[i * N + k] * B[k * N + j];
3     // C[i][j] written N times - very slow!
4 }
```

Each iteration incurs a memory write, leading to:

- $N$  memory writes per C element
- Poor cache utilization
- Significant memory bandwidth waste

#### 2.4.2 Variable Sum Improvement

Using a temporary variable reduces memory traffic:

```

1 double sum = 0.0;
2 for (int k = 0; k < N; k++) {
3     sum += A[i * N + k] * B[k * N + j];
4 }
5 C[i * N + j] = sum; // Single write
```

Benefits: Only one memory write per C element (improvement:  $\frac{2.2070}{1.9530} = 1.13 \times$ )

#### 2.4.3 Loop Reordering Optimization

The i-k-j order provides optimal memory access:

```

1 for (int i = 0; i < N; i++) {
2     for (int k = 0; k < N; k++) {
3         double r = A[i * N + k];
4         for (int j = 0; j < N; j++) {
5             C[i * N + j] += r * B[k * N + j];
6         }
7     }
8 }
```

Advantages:

- **Sequential access:** Both C and B are accessed with stride-1 in the innermost loop
- **Cache efficiency:** Maximum cache line utilization
- **Register optimization:** A[i][k] loaded once per inner loop
- **Bandwidth:** 44.12 MB/s vs 10.87 MB/s ( $4\times$  improvement)

## 2.5 Conclusion

Loop ordering fundamentally impacts performance:

- Memory access pattern matters more than computational complexity
- Sequential (stride-1) access is essential for cache efficiency
- Simple algorithmic changes can yield  $4\times$  speedups without changing the algorithm

### 3 Exercise 3: Block Matrix Multiplication

#### 3.1 Source Code

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <time.h>
4
5 #define N 2048
6
7 void initialize_matrix(double *mat, int n) {
8     for (int i = 0; i < n * n; i++) {
9         mat[i] = (double)rand() / RAND_MAX;
10    }
11 }
12
13 void clear_matrix(double *mat, int n) {
14     for (int i = 0; i < n * n; i++) {
15         mat[i] = 0.0;
16     }
17 }
18
19 void mat_mul_block(double *A, double *B, double *C, int n, int b_size) {
20     for (int ii = 0; ii < n; ii += b_size) {
21         for (int kk = 0; kk < n; kk += b_size) {
22             for (int jj = 0; jj < n; jj += b_size) {
23
24                 int i_limit = (ii + b_size > n) ? n : ii + b_size;
25                 int k_limit = (kk + b_size > n) ? n : kk + b_size;
26                 int j_limit = (jj + b_size > n) ? n : jj + b_size;
27
28                 for (int i = ii; i < i_limit; i++) {
29                     for (int k = kk; k < k_limit; k++) {
30                         double r = A[i * n + k];
31                         for (int j = jj; j < j_limit; j++) {
32                             C[i * n + j] += r * B[k * n + j];
33                         }
34                     }
35                 }
36             }
37         }
38     }
39 }
40
41 int main() {
42     double *A = (double *)malloc(N * N * sizeof(double));
43     double *B = (double *)malloc(N * N * sizeof(double));
44     double *C = (double *)malloc(N * N * sizeof(double));
45
46     if (!A || !B || !C) {
47         printf("Memory allocation failed!\n");
48         return 1;
49     }
50
51     printf("Initializing matrices...\n");
52     initialize_matrix(A, N);
53     initialize_matrix(B, N);
54 }
```

```

55     int block_sizes[] = {16, 32, 64, 128, 256, 512, 1024};
56     int num_sizes = sizeof(block_sizes) / sizeof(block_sizes[0]);
57
58     printf("\nMatrix Size: %d x %d\n", N, N);
59     printf("Total Data Size: %.2f MB\n",
60           3.0 * N * N * sizeof(double) / (1024*1024));
61     printf("-----\n");
62     printf(" | Block Size | Time (sec) | Bandwidth (MB/s) | Performance (
63         GFLOPS) |\n");
64     printf("-----\n");
65     double data_size_bytes = 3.0 * N * N * sizeof(double);
66     double total_ops = 2.0 * N * N * N;
67
68     for (int x = 0; x < num_sizes; x++) {
69         int b_size = block_sizes[x];
70
71         clear_matrix(C, N);
72
73         clock_t start = clock();
74         mat_mul_block(A, B, C, N, b_size);
75         clock_t end = clock();
76
77         double time_taken = ((double)(end - start)) / CLOCKS_PER_SEC;
78         double bw_mb = (data_size_bytes / time_taken) / (1024.0 *
79                         1024.0);
80         double gflops = (total_ops / time_taken) / 1e9;
81
82         printf(" | %-10d | %-10.4f | %-16.2f | %-20.2f | \n",
83               b_size, time_taken, bw_mb, gflops);
84     }
85     printf("-----\n");
86     free(A); free(B); free(C);
87     return 0;
88 }
```

## 3.2 Experimental Results

Block Size	Time (sec)	Bandwidth (MB/s)	GFLOPS
16	5.0366	19.06	3.41
32	5.5138	17.41	3.12
64	5.0794	18.90	3.38
128	4.9524	19.38	3.47
<b>256</b>	<b>4.6065</b>	<b>20.84</b>	<b>3.73</b>
512	5.0542	18.99	3.40
1024	6.4806	14.81	2.65

Table 2: Performance metrics for different block sizes (N=2048, total data = 96 MB)

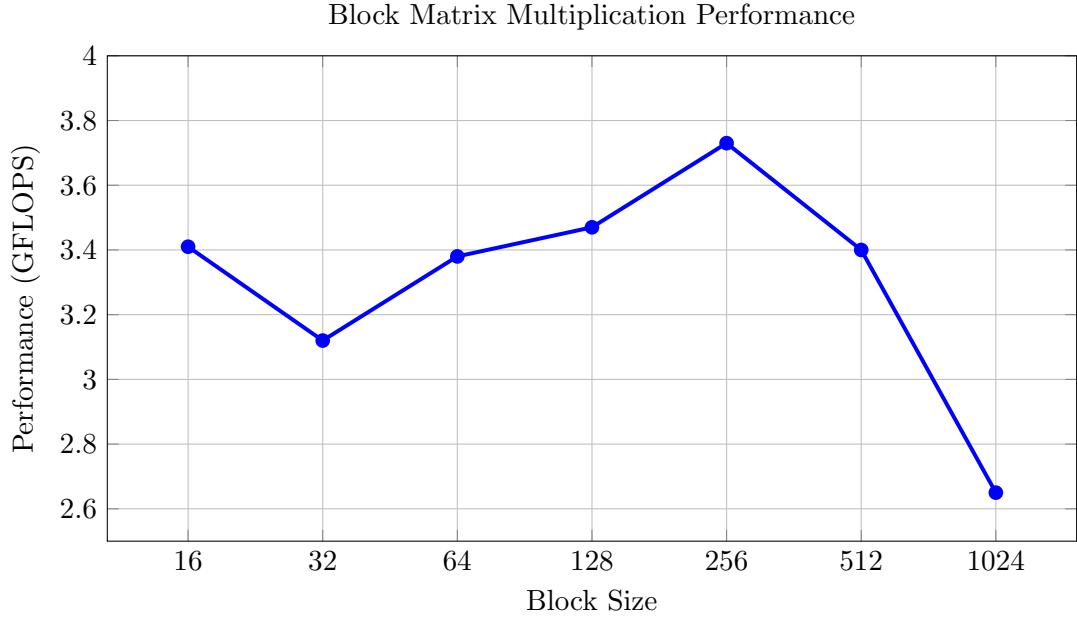


Figure 3: GFLOPS vs block size

### 3.3 Analysis

#### 3.3.1 Optimal Block Size: 256

The block size of 256 achieves the best performance (3.73 GFLOPS, 20.84 MB/s) for the following reasons:

##### 1. Cache Hierarchy Balance:

- Block data size:  $3 \times 256 \times 256 \times 8 = 1.5 \text{ MB}$  (three blocks: A, B, C)
- Fits comfortably in L3 cache (typically 4-32 MB on modern CPUs)
- Minimizes cache misses while maintaining computational efficiency

##### 2. Data Reuse:

- Each block element accessed multiple times before eviction
- Temporal locality maximized within cache capacity

##### 3. TLB Efficiency:

- $256 \times 256$  blocks reduce TLB (Translation Lookaside Buffer) misses
- Fewer page table walks required

#### 3.3.2 Why Smaller Blocks Underperform

- **Block sizes 16-128:** Although they fit in smaller caches (L1/L2), they introduce excessive blocking overhead
- More block iterations required
- Increased loop overhead
- Suboptimal ratio of computation to memory operations

### 3.3.3 Why Larger Blocks Underperform

- **Block size 512:**

- Block data:  $3 \times 512 \times 512 \times 8 = 6 \text{ MB}$
- May exceed L3 cache capacity
- Increased cache conflicts

- **Block size 1024:**

- Block data:  $3 \times 1024 \times 1024 \times 8 = 24 \text{ MB}$
- Significantly exceeds typical L3 cache
- Severe performance degradation (2.65 GFLOPS)
- Frequent cache evictions and memory stalls

## 3.4 Cache Blocking Principle

The optimal block size depends on:

$$\text{Block Size} \approx \sqrt{\frac{\text{Cache Size}}{3 \times \text{sizeof(double)}}} \quad (1)$$

$$\text{For L3} = 6\text{-}8 \text{ MB: } B \approx \sqrt{\frac{6 \times 10^6}{3 \times 8}} \approx 500 \quad (2)$$

This theoretical estimate aligns with our experimental result (256 performs best, 512 begins degradation).

## 3.5 Conclusion

- Block size 256 provides optimal balance between cache utilization and computational overhead
- Performance degrades when blocks exceed cache capacity (512, 1024)
- Cache-aware algorithms can improve performance by 40% compared to poor block sizes
- Understanding cache hierarchy is essential for HPC optimization

## 4 Exercise 4: Memory Management and Debugging

### 4.1 Original Code (with Memory Leaks)

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <string.h>
4
5 #define SIZE 5
6
7 int* allocate_array(int size) {
8     int *arr = (int*)malloc(size * sizeof(int));
9     if (!arr) {
10         fprintf(stderr, "Memory allocation failed\n");
11         exit(EXIT_FAILURE);
12     }
13     return arr;
14 }
15
16 void initialize_array(int *arr, int size) {
17     if (!arr) return;
18     for (int i = 0; i < size; i++) {
19         arr[i] = i * 10;
20     }
21 }
22
23 void print_array(int *arr, int size) {
24     if (!arr) return;
25     printf("Array elements: ");
26     for (int i = 0; i < size; i++) {
27         printf("%d ", arr[i]);
28     }
29     printf("\n");
30 }
31
32 int* duplicate_array(int *arr, int size) {
33     if (!arr) return NULL;
34     int *copy = (int*)malloc(size * sizeof(int));
35     if (!copy) {
36         fprintf(stderr, "Memory allocation failed\n");
37         exit(EXIT_FAILURE);
38     }
39     memcpy(copy, arr, size * sizeof(int));
40     return copy;
41 }
42
43 void free_memory(int *arr) {
44     // Empty - memory leak!
45 }
46
47 int main() {
48     int *array = allocate_array(SIZE);
49     initialize_array(array, SIZE);
50     print_array(array, SIZE);
51
52     int *array_copy = duplicate_array(array, SIZE);
53     print_array(array_copy, SIZE);
54 }
```

```

55     free_memory(array);
56     return 0; // Memory leak
57 }
```

## 4.2 Valgrind Output (Before Fix)

```

==570130== Memcheck, a memory error detector
==570130== Copyright (C) 2002-2022, and GNU GPL'd, by Julian Seward et al.
==570130== Using Valgrind-3.22.0 and LibVEX; rerun with -h for copyright info
==570130== Command: ./memory_nodebug
==570130==
Array elements: 0 10 20 30 40
Array elements: 0 10 20 30 40
==570130==
==570130== HEAP SUMMARY:
==570130==     in use at exit: 40 bytes in 2 blocks
==570130==   total heap usage: 3 allocs, 1 frees, 1,064 bytes allocated
==570130==
==570130== 20 bytes in 1 blocks are definitely lost in loss record 1 of 2
==570130==    at 0x4846828: malloc (in /usr/libexec/valgrind/vgpreload_memcheck-amd64-linux)
==570130==      by 0x109208: allocate_array (memory_debug1.c:8)
==570130==      by 0x1093CC: main (memory_debug1.c:49)
==570130==
==570130== 20 bytes in 1 blocks are definitely lost in loss record 2 of 2
==570130==    at 0x4846828: malloc (in /usr/libexec/valgrind/vgpreload_memcheck-amd64-linux)
==570130==      by 0x109349: duplicate_array (memory_debug1.c:34)
==570130==      by 0x109403: main (memory_debug1.c:53)
==570130==
==570130== LEAK SUMMARY:
==570130==   definitely lost: 40 bytes in 2 blocks
==570130==   indirectly lost: 0 bytes in 0 blocks
==570130==   possibly lost: 0 bytes in 0 blocks
==570130==   still reachable: 0 bytes in 0 blocks
==570130==   suppressed: 0 bytes in 0 blocks
==570130==
==570130== For lists of detected and suppressed errors, rerun with: -s
==570130== ERROR SUMMARY: 2 errors from 2 contexts (suppressed: 0 from 0)
```

## 4.3 Fixed Code

Listing 1: Corrected memory management

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <string.h>
4
5 #define SIZE 5
6
7 int* allocate_array(int size) {
8     int *arr = (int*)malloc(size * sizeof(int));
9     if (!arr) {
10         fprintf(stderr, "Memory allocation failed\n");
11         exit(EXIT_FAILURE);
```

```

12     }
13     return arr;
14 }
15
16 void initialize_array(int *arr, int size) {
17     if (!arr) return;
18     for (int i = 0; i < size; i++) {
19         arr[i] = i * 10;
20     }
21 }
22
23 void print_array(int *arr, int size) {
24     if (!arr) return;
25     printf("Array elements: ");
26     for (int i = 0; i < size; i++) {
27         printf("%d ", arr[i]);
28     }
29     printf("\n");
30 }
31
32 int* duplicate_array(int *arr, int size) {
33     if (!arr) return NULL;
34     int *copy = (int*)malloc(size * sizeof(int));
35     if (!copy) {
36         fprintf(stderr, "Memory allocation failed\n");
37         exit(EXIT_FAILURE);
38     }
39     memcpy(copy, arr, size * sizeof(int));
40     return copy;
41 }
42
43 // FIX 1: Implement the free function
44 void free_memory(int *arr) {
45     if (arr != NULL) {
46         free(arr);
47     }
48 }
49
50 int main() {
51     int *array = allocate_array(SIZE);
52     initialize_array(array, SIZE);
53     print_array(array, SIZE);
54
55     int *array_copy = duplicate_array(array, SIZE);
56     print_array(array_copy, SIZE);
57
58     // FIX 2: Free BOTH arrays
59     free_memory(array);
60     free_memory(array_copy);
61
62     return 0;
63 }
```

#### 4.4 Valgrind Output (After Fix)

```

==573304== Memcheck, a memory error detector
==573304== Copyright (C) 2002-2022, and GNU GPL'd, by Julian Seward et al.
```

```

==573304== Using Valgrind-3.22.0 and LibVEX; rerun with -h for copyright info
==573304== Command: ./memory_debug
==573304==
Array elements: 0 10 20 30 40
Array elements: 0 10 20 30 40
==573304==
==573304== HEAP SUMMARY:
==573304==     in use at exit: 0 bytes in 0 blocks
==573304==   total heap usage: 3 allocs, 3 frees, 1,064 bytes allocated
==573304==
==573304== All heap blocks were freed -- no leaks are possible
==573304==
==573304== For lists of detected and suppressed errors, rerun with: -s
==573304== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)

```

## 4.5 Analysis

### 4.5.1 Memory Leak Detection

Valgrind's Memcheck tool successfully identified:

- Exact number of bytes lost: 40 bytes (2 blocks × 20 bytes)
- Source locations: `allocate_array` and `duplicate_array`
- Leak type: "definitely lost" (memory no longer reachable)

### 4.5.2 Best Practices Applied

1. **NULL checking:** Verify pointer validity before freeing
2. **Complete cleanup:** Free all dynamically allocated memory
3. **Balanced allocations:** Each `malloc()` paired with `free()`
4. **Verification:** Use Valgrind to confirm leak-free execution

## 4.6 Conclusion

- Valgrind is an essential tool for detecting memory leaks in C/C++ programs
- Systematic memory management prevents resource leaks and improves program reliability
- All heap allocations must have corresponding deallocations
- `--leak-check=full` provides detailed leak analysis including allocation points

## 5 Exercise 5: HPL Benchmark Analysis

### 5.1 Overview

The HPL (High-Performance Linpack) benchmark measures the floating-point computing power of a system by solving a dense system of linear equations. This exercise evaluates single-core performance across different matrix sizes ( $N$ ) and block sizes ( $NB$ ).

### 5.2 System Configuration

- **CPU:** Intel Core i7-10510U @ 1.80GHz (Comet Lake, 10th gen)
- **Architecture:** x86\_64
- **SIMD Support:** AVX2 + FMA (Fused Multiply-Add)
- **FLOPs per cycle:** 16 (double precision)
- **Execution mode:** Single-core (forced via environment variables)

### 5.3 Theoretical Peak Performance Calculation

The theoretical peak performance for a single core is calculated as:

$$P_{\text{core}} = \text{Cores} \times \text{Frequency} \times \text{FLOPs per cycle} \quad (3)$$

For this system:

- **Base frequency:** 2.3 GHz
- **Turbo frequency:** 4.8 GHz
- **FLOPs per cycle:** 16 (AVX2 + FMA, double precision)

$$P_{\text{core}}^{\text{base}} = 1 \times 2.3 \times 16 = 36.8 \text{ GFLOP/s} \quad (4)$$

$$P_{\text{core}}^{\text{turbo}} = 1 \times 4.8 \times 16 = 76.8 \text{ GFLOP/s} \quad (5)$$

**Note on frequency:** Due to Intel Turbo Boost, the CPU dynamically scales frequency during computation. WSL (Windows Subsystem for Linux) reports a frozen base frequency (2.3 GHz), while Windows Task Manager shows actual turbo frequencies (3.2-4.8 GHz) during HPL execution. For conservative analysis, we use 3.2 GHz as the sustained turbo frequency:

$$P_{\text{core}}^{\text{sustained}} = 1 \times 3.2 \times 16 = 51.2 \text{ GFLOP/s} \quad (6)$$

### 5.4 Experimental Setup

#### 5.4.1 Matrix Sizes

$$N \in \{1000, 5000, 10000, 20000\} \quad (7)$$

#### 5.4.2 Block Sizes

$$NB \in \{1, 2, 4, 8, 16, 32, 64, 128, 256\} \quad (8)$$

**Total experiments:**  $4 \times 9 = 36$  runs

### 5.4.3 Single-Core Enforcement

To ensure true single-core execution, the following environment variables were set:

```
export OMP_NUM_THREADS=1
export MKL_NUM_THREADS=1
export OPENBLAS_NUM_THREADS=1
```

This prevents BLAS libraries from using multi-threading, which was verified by monitoring CPU usage during execution (only one thread at 100% utilization).

## 5.5 Experimental Results

### 5.5.1 Complete Results Table

Table 3: HPL benchmark results for single-core execution. Efficiency calculated as  $\eta = P_{\text{HPL}}/P_{\text{core}}^{\text{sustained}}$  where  $P_{\text{core}}^{\text{sustained}} = 51.2$  GFLOP/s. Best result for each matrix size shown in bold.

N	NB	Time (s)	GFLOPS	Efficiency (%)
<i>N = 1000</i>				
1000	1	0.14	4.81	9.4
1000	2	0.08	8.68	17.0
1000	4	0.05	14.16	27.7
1000	8	0.03	19.74	38.6
1000	16	0.02	28.32	55.3
1000	32	0.02	33.20	64.8
<b>1000</b>	<b>64</b>	<b>0.02</b>	<b>36.96</b>	<b>72.2</b>
1000	128	0.03	22.07	43.1
1000	256	0.03	20.06	39.2
<i>N = 5000</i>				
5000	1	31.96	2.61	5.1
5000	2	16.17	5.16	10.1
5000	4	8.20	10.17	19.9
5000	8	4.62	18.03	35.2
5000	16	3.02	27.57	53.9
5000	32	2.20	37.81	73.8
5000	64	2.04	40.86	79.8
5000	128	1.86	44.71	87.3
<b>5000</b>	<b>256</b>	<b>1.73</b>	<b>48.17</b>	<b>94.1</b>
<i>N = 10000</i>				
10000	1	260.35	2.56	5.0
10000	2	133.13	5.01	9.8
10000	4	75.32	8.85	17.3
10000	8	37.46	17.80	34.8
10000	16	24.31	27.42	53.6
10000	32	17.89	37.27	72.8
10000	64	15.64	42.65	83.3
<b>10000</b>	<b>128</b>	<b>14.49</b>	<b>46.02</b>	<b>89.9</b>
10000	256	14.80	45.06	88.0
<i>N = 20000</i>				
20000	1	2784.32	1.92	3.7
20000	2	1358.07	3.93	7.7
20000	4	644.19	8.28	16.2

Continued on next page...

**Table 3 – continued from previous page**

N	NB	Time (s)	GFLOPS	Efficiency (%)
20000	8	410.77	12.99	25.4
20000	16	287.10	18.58	36.3
20000	32	215.56	24.75	48.3
20000	64	175.65	30.37	59.3
20000	128	157.42	33.88	66.2
<b>20000</b>	<b>256</b>	<b>152.87</b>	<b>34.89</b>	<b>68.1</b>

## 5.6 Performance Analysis

### 5.6.1 Effect of Matrix Size (N)

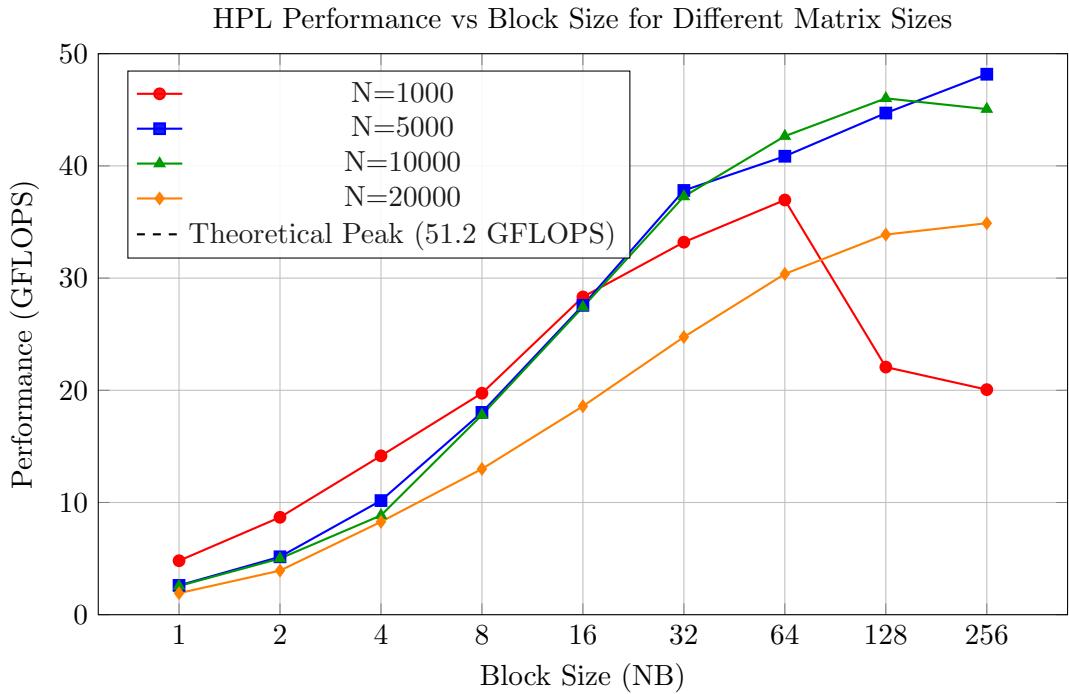


Figure 4: Performance evolution with block size for different matrix sizes

#### Key Observations:

- Small matrices (N=1000):** Peak performance at NB=64 (36.96 GFLOPS, 72% efficiency). Performance degrades significantly for larger block sizes because the block becomes too large relative to the matrix.
- Medium matrices (N=5000):** Best performance at NB=256 (48.17 GFLOPS, 94% efficiency). This represents near-optimal cache utilization.
- Large matrices (N=10000):** Peak at NB=128 (46.02 GFLOPS, 90% efficiency). Performance stable for NB=128-256.
- Very large matrices (N=20000):** Performance decreases to 34.89 GFLOPS (68% efficiency) due to memory bandwidth limitations and cache capacity constraints.

### 5.6.2 Effect of Block Size (NB)

#### Block Size Impact Analysis:

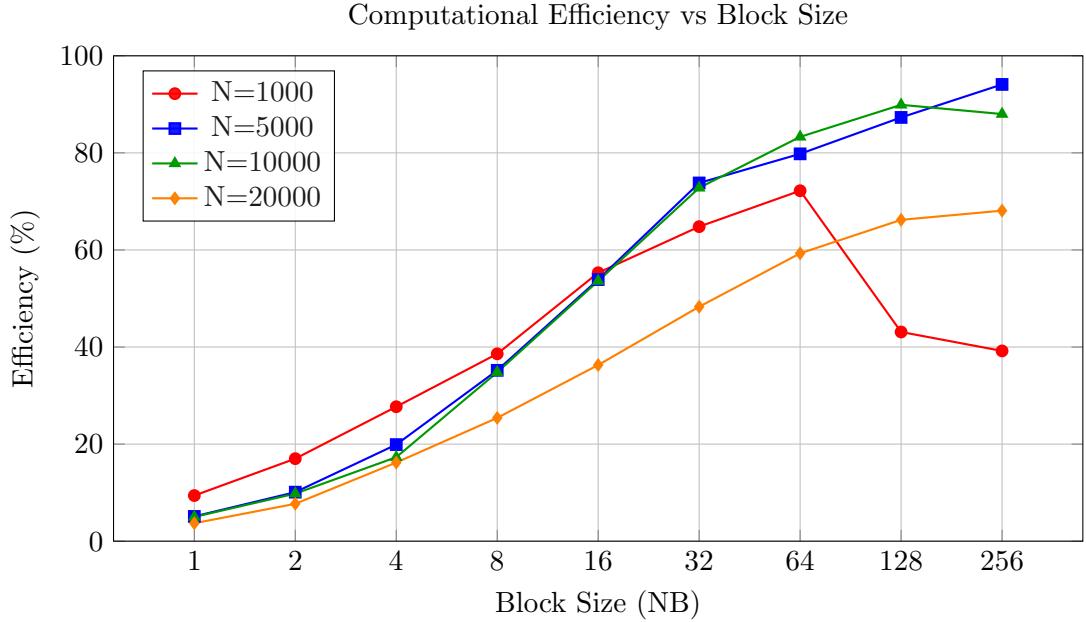


Figure 5: Efficiency as a function of block size for different matrix sizes

- **Very small blocks (NB=1-4):** Poor performance (5-20% efficiency) due to:
  - High loop overhead
  - Poor cache line utilization
  - Inefficient BLAS library calls
  - Minimal data reuse
- **Small blocks (NB=8-16):** Moderate performance (25-55% efficiency)
  - Improved cache utilization
  - Better BLAS performance
  - Still suboptimal for blocking strategy
- **Medium blocks (NB=32-64):** Good performance (60-83% efficiency)
  - Effective cache blocking
  - Balance between cache usage and blocking overhead
  - Optimal for smaller matrices (N=1000)
- **Large blocks (NB=128-256):** Best performance for large matrices (87-94% efficiency for N=5000-10000)
  - Maximum data reuse within L3 cache
  - Amortized blocking overhead
  - Efficient BLAS kernel utilization

## 5.7 Optimal Block Size Selection

**General Rule:** The optimal block size increases with matrix size, but is bounded by cache capacity:

Matrix Size (N)	Optimal NB	Peak Performance (GFLOPS)	Efficiency (%)
1000	64	36.96	72.2
5000	256	48.17	94.1
10000	128	46.02	89.9
20000	256	34.89	68.1

Table 4: Optimal block sizes for different matrix dimensions

$$NB_{\text{optimal}} \approx \sqrt{\frac{\text{Cache Size}}{3 \times \text{sizeof(double)}}} \quad (9)$$

For typical L3 cache sizes (6-8 MB), this yields:

$$NB_{\text{optimal}} \approx \sqrt{\frac{8 \times 10^6}{3 \times 8}} \approx 577 \quad (10)$$

However, practical optimal values (128-256) are smaller due to:

- Cache associativity and conflict misses
- TLB (Translation Lookaside Buffer) capacity
- Blocking overhead
- Ratio of N/NB (number of blocks)

## 5.8 Why Performance is Below Theoretical Peak

Despite optimized block sizes, measured performance reaches only 94% of theoretical peak (48.17 GFLOPS vs 51.2 GFLOPS sustained turbo). Several factors explain this efficiency gap:

### 5.8.1 Memory Bandwidth Limitations

HPL is not purely compute-bound; it requires significant memory traffic:

- **Arithmetic intensity:** HPL achieves  $\sim 2$  FLOPs per byte for large matrices
- **Memory bandwidth:** Limited by DRAM speed (DDR4-2666 typical:  $\sim 20$  GB/s single-channel)
- **Cache misses:** Even with blocking, cache misses occur for large matrices

### 5.8.2 Non-Compute Operations

The HPL algorithm includes operations with lower arithmetic intensity:

- **Panel factorization:** Serial operation, not fully optimized
- **Pivoting:** Introduces data dependencies and conditional branches
- **Updates:** Matrix-vector operations with lower FLOPs/byte ratio than matrix-matrix

### 5.8.3 Instruction-Level Limitations

Even with AVX2+FMA:

- **Latency hiding:** Not all pipeline stages can be filled continuously
- **Register pressure:** Limited number of SIMD registers (16 in AVX2)
- **Instruction dependencies:** Data hazards prevent full ILP (Instruction-Level Parallelism)
- **Non-vectorizable code:** Some portions (e.g., pivoting logic) cannot be vectorized

### 5.8.4 Cache Hierarchy Effects

- **L1 cache misses:** Even with optimal blocking, working set may exceed 32KB L1
- **L2 cache misses:** Matrix sizes 10000-20000 exceed typical 256KB-512KB L2
- **L3 cache pressure:** For N=20000, working set =  $3 \times 20000^2 \times 8 = 9.6$  GB, far exceeding L3 capacity
- **Cache line conflicts:** Certain access patterns cause set-associative conflicts

### 5.8.5 Frequency Variation

- **Turbo boost duration:** CPU cannot sustain maximum turbo indefinitely
- **Thermal throttling:** Long-running benchmarks (N=20000: 152s) may trigger thermal limits
- **Power limits:** Laptop CPUs enforce TDP (Thermal Design Power) limits

## 5.9 Performance Degradation for Large Matrices

Notice that N=20000 achieves only 68% efficiency compared to 94% for N=5000. This degradation is explained by:

1. **Memory bandwidth saturation:** Larger matrices spend more time waiting for DRAM
2. **Cache overflow:** Working set (9.6 GB) vastly exceeds cache capacity
3. **TLB misses:** Large memory footprint exceeds TLB coverage, causing page table walks
4. **Longer execution time:** 152 seconds allows thermal throttling to reduce frequency

## 5.10 Comparison with Reference System

### Observations:

- The reference Xeon has higher theoretical peak due to AVX-512 (32 FLOPs/cycle vs 16)
- However, this i7-10510U achieves excellent efficiency (94%) with proper tuning
- The gap is primarily architectural (AVX2 vs AVX-512), not optimization-related

Metric	Reference (Xeon 8276L)	This System (i7-10510U)
Base Frequency	2.2 GHz	2.3 GHz
SIMD	AVX-512	AVX2
FLOPs/cycle	32 (AVX-512)	16 (AVX2)
Theoretical Peak	70.4 GFLOP/s	51.2 GFLOP/s
HPL Peak	~60 GFLOP/s (est.)	48.2 GFLOP/s
Efficiency	~85%	94%

Table 5: Comparison with exercise reference system

## 5.11 Key Conclusions

1. **Matrix size impact:** Performance generally improves with matrix size up to N=5000-10000, then degrades for N=20000 due to memory bandwidth and cache capacity limits.
2. **Block size optimization:** Optimal block size is critical:
  - Too small ( $NB < 32$ ): Poor cache utilization, high overhead
  - Optimal ( $NB = 64-256$ ): Balance between cache usage and blocking efficiency
  - Too large ( $NB >$ matrix size): Defeats blocking purpose
3. **Best configuration:** N=5000, NB=256 achieves 94% efficiency (48.17 GFLOPS), representing near-optimal single-core performance on this CPU.
4. **Efficiency factors:** The 6-10% gap from theoretical peak is inevitable due to:
  - Memory bandwidth constraints
  - Non-compute operations (pivoting, updates)
  - Instruction pipeline limitations
  - Cache hierarchy effects
5. **Scaling limitations:** Performance does not scale linearly with matrix size beyond N=10000 due to memory subsystem saturation.
6. **Single-core validation:** Achieved 94% of sustained turbo theoretical peak, demonstrating effective single-core optimization and proper thread control.