

# TP3 - OpenMP Analysis

## Understanding Parallel Computing Fundamentals

Kawtar Labzae

February 2026

### 1 Exercise 1: Fork-Join Model and Thread Management

#### 1.1 Essential Parallelization Code

```
1 #pragma omp parallel private(rank)
2 {
3     rank = omp_get_thread_num();
4     printf("Hello from rank %d\n", rank);
5     if (rank == 0) // Only master thread
6         nthreads = omp_get_num_threads();
7 }
```

#### 1.2 Conceptual Framework

**Fork-Join Execution Model:**

1. **Sequential region:** Single master thread
2. **Fork:** Master spawns worker threads at parallel directive
3. **Parallel region:** All threads execute concurrently
4. **Join:** Implicit barrier, workers terminate, master continues

**Why Restrict Thread 0?** The assignment `nthreads = ...` creates a write operation. If all N threads write simultaneously:

- Creates unnecessary memory bus contention
- Triggers cache coherence protocol overhead
- Violates single-responsibility principle

**Non-deterministic Execution:** Thread output order varies because OS scheduler controls:

- Core assignment

- Context switching timing
- Priority levels

**Critical Design Principle:** Parallel algorithms must be **order-independent**. Any dependency on execution order creates race conditions.

## 2 Exercise 2: Reduction Patterns and Work Distribution

### 2.1 Core Challenge

Computing  $\pi \approx \sum_{i=0}^{n-1} \frac{4}{1+x_i^2} \Delta x$  requires accumulating partial sums without race conditions.

```

1 #pragma omp parallel reduction(:global_sum)
2 {
3     id = omp_get_thread_num();
4     nthreads = omp_get_num_threads();
5     // Cyclic distribution
6     for (i = id; i < num_steps; i += nthreads) {
7         x = (i + 0.5) * step;
8         global_sum += 4.0 / (1.0 + x * x);
9     }
10 }
```

### 2.2 Reduction Mechanism Internals

The Race Condition Without Reduction:

Time	Thread 0	Thread 1	Memory(sum)
t0	read sum=0	read sum=0	0
t1	compute +5	compute +3	0
t2	write sum=5	write sum=3	5 (or 3!)

Result: Lost update! Final sum is wrong.

How reduction(:var) Solves This:

1. **Privatization:** Creates `sum_private[nthreads]`
2. **Initialization:** Each `sum_private[i] = 0` (identity for +)
3. **Local accumulation:** Thread i updates only `sum_private[i]`
4. **Barrier:** Implicit synchronization
5. **Final reduction:** `sum = sum_private[0] + sum_private[1] + ...`

## 2.3 Work Distribution Strategies

Strategy	Iteration Assignment	Memory Access
Cyclic (used)	T0:0,4,8... T1:1,5,9...	Identical (Compute-bound)
Block	T0:0-24999, T1:25000-49999	Identical (Compute-bound)

**Correction on Cache:** Previous assumptions about stride are incorrect here. Since  $x$  is calculated mathematically ( $x = (i + 0.5) * step$ ) rather than loaded from an array, there is no memory stride penalty. Both strategies have identical cache behavior.

**Why Cyclic Here?** Demonstrates manual control. It ensures load balancing if iterations had variable costs, but for this uniform loop, it adds complexity.

## 3 Exercise 3: Abstraction vs Manual Control

### 3.1 The Minimal Parallelization

```
1 // Before: Serial
2 for (i = 0; i < num_steps; i++) {
3     sum += 4.0 / (1.0 + x * x);
4 }
5
6 // After: One-line addition
7 #pragma omp parallel for reduction(+:sum) private(x)
8 for (i = 0; i < num_steps; i++) {
9     sum += 4.0 / (1.0 + x * x);
10 }
```

### 3.2 What OpenMP Automates

Aspect	Automatic Handling by <code>parallel for</code>
Thread creation	Fork-join managed internally
Work distribution	Default: static block scheduling
Variable scoping	Loop index <code>i</code> automatically private
Synchronization	Implicit barrier before reduction

### 3.3 Performance Comparison

	Ex2 (Manual)	Ex3 (Auto)
Time	0.001373s	0.002336s
Distribution	Cyclic	Block

**Why is Auto Slower? (Runtime Overhead)** The manual version is faster at this small scale because it avoids the abstraction tax of OpenMP.

1. **Library Calls:** `pragma omp for` invokes runtime library functions to calculate loop bounds.
2. **Safety Checks:** The runtime manages implicit barriers and safety checks that the manual `int i = id` logic skips.
3. **Scale:** At 100k iterations, the computation is so fast that this runtime overhead (0.5-1ms) becomes a significant percentage of total time.

**Practical Lesson:** `parallel for` is safer and cleaner, but for extremely small loops, manual distribution can be slightly faster by bypassing runtime management.

## 4 Exercise 4: Matrix Multiplication - Memory Bounds and Scheduling

### 4.1 Parallelization Structure

```

1 #pragma omp parallel for collapse(2) schedule(runtime)
2 for (int i = 0; i < 1000; i++) {
3     for (int j = 0; j < 1000; j++) {
4         for (int k = 0; k < 1000; k++) {
5             c[i*m + j] += a[i*n + k] * b[k*m + j];
6         }
    }}
```

### 4.2 Why `collapse(2)`?

**Without collapse:**

- Parallelizes only i-loop: 1000 iterations
- With 16 threads: 62 iterations/thread

**With collapse(2):**

- Merges i and j:  $1000 \times 1000 = 1,000,000$  iterations
- With 16 threads: 62,500 iterations/thread

**Benefit:** Finer granularity → better load balance when threads iterations.

**Why NOT collapse(3)?** The k-loop accumulates into `c[i][j]`. Multiple threads would write to same location → requires synchronization → kills performance.

## 4.3 Experimental Results Analysis

Threads	Sched	Chunk	Time (s)	Speedup
1	static	100	7.640	1.00
2	static	100	3.373	2.27
4	guided	100	2.620	2.92
8	dynamic	10	2.453	3.11
16	guided	1	2.748	2.78

Table 1: Best configuration for each thread count

## 4.4 Critical Performance Insights

### 4.4.1 Why Only 3.11x Speedup with 8 Threads?

#### 1. Memory Bandwidth Saturation

Matrix size:  $3 \times 1000^2 \times 8$  bytes = 24 MB

Each element of C requires:

- 1000 reads from A (row i)
- 1000 reads from B (column j)
- 1 write to C

Total memory traffic:  $1000^3 \times (2 \times 8 \text{ bytes read} + 8 \text{ write}) = 16 \text{ GB}$

**Typical memory bandwidth:** 20-40 GB/s per socket

**Implication:** Adding threads doesn't help when all threads are waiting for memory!

#### 2. Cache Inefficiency

- Matrix B accessed column-wise:  $b[k][j]$
- Non-contiguous access → new cache line every iteration
- L1 cache ( 32 KB)  $\ll$  column size (8000 bytes)
- Cache miss rate  $\approx$  90-95%

#### 3. False Sharing

Adjacent elements of C (e.g.,  $c[i][j]$  and  $c[i][j+1]$ ) often share a 64-byte cache line. When different threads write nearby elements:

- Each write invalidates other cores' caches
- Cache coherence protocol overhead
- Ping-ponging of cache lines between cores

## 4.5 Scheduling Strategy Analysis

Schedule	Characteristics
static	Chunks assigned at compile/start time. <b>Lowest overhead</b> , best for uniform work.
dynamic	Runtime work queue. Thread finishes chunk → gets next from queue. <b>Higher overhead</b> but handles load imbalance.
guided	Starts with large chunks, decreases exponentially. Compromise between static and dynamic.

Why dynamic(10) wins at 8 threads?

- Matrix multiplication has relatively uniform work per iteration
- But cache effects create some variance
- Dynamic handles variance without excessive overhead
- Chunk=10 balances scheduling cost vs granularity

## 4.6 The 16-Thread Performance Drop

Speedup at 8 threads: 3.11x

Speedup at 16 threads: 2.78x (slower!)

**Root Causes:**

1. **Memory bandwidth fully saturated:** All 16 threads fighting for same bus
2. **Thread overhead:** Context switching, synchronization costs scale with thread count

## 4.7 Key Lessons

1. **Identify the bottleneck:** CPU-bound vs memory-bound vs synchronization-bound
2. **Memory bandwidth is finite:** Adding threads to memory-bound code doesn't help
3. **More threads ≠ faster:** Overhead and contention can reduce performance

# 5 Exercise 5: Parallelizing Jacobi Method with OpenMP

## 5.1 Code Implementation

The Jacobi method solves a system of linear equations iteratively. The algorithm contains two main computational intense loops that were parallelized:

1. Computing the new vector approximation ( $x_{courant}$ ).
2. Checking for convergence (finding the maximum error).

Below are the relevant modified sections of the code using OpenMP directives:

Listing 1: Parallelized Jacobi Iteration

```

1  while (1) {
2      iteration++;
3
4      // Parallelize the calculation of x_courant
5      // 'j' must be private to avoid race conditions in the inner
6      // loops
7      #pragma omp parallel for private(j)
8      for (i = 0; i < n; i++) {
9          x_courant[i] = 0;
10         for (j = 0; j < i; j++) {
11             x_courant[i] += a[j * n + i] * x[j];
12         }
13         for (j = i + 1; j < n; j++) {
14             x_courant[i] += a[j * n + i] * x[j];
15         }
16         x_courant[i] = (b[i] - x_courant[i]) / a[i * n + i];
17     }
18
19     double absmax = 0;
20
21     // Parallelize the convergence check
22     // Use reduction to safely find the global maximum error
23     #pragma omp parallel for reduction(max: absmax)
24     for (i = 0; i < n; i++) {
25         double curr = fabs(x[i] - x_courant[i]);
26         if (curr > absmax)
27             absmax = curr;
28     }
29
30     norme = absmax / n;
31     // ... convergence check and swap ...
}

```

## 5.2 Performance Analysis

The code was executed with  $N = 2000$  iterations using 1, 2, 4, 8, and 16 threads.

### 5.2.1 Experimental Data

$$\text{Speedup}(S_p) = \frac{T_1}{T_p} \quad , \quad \text{Efficiency}(E_p) = \frac{S_p}{p}$$

Threads (p)	Time (s)	Speedup ( $S_p$ )	Efficiency ( $E_p$ )
1	58.43	1.00	1.00 (100%)
2	36.22	1.61	0.81 (81%)
4	26.74	2.19	0.55 (55%)
8	23.67	2.47	0.31 (31%)
16	25.69	2.27	0.14 (14%)

Table 2: Performance Metrics for Jacobi Method ( $N = 2000$ )

### 5.2.2 Speedup and Efficiency Plots




Figure 1: The speedup plateaus significantly after 4 threads and degrades at 16.




Figure 2: Efficiency drops rapidly as thread count increases.

### 5.3 Discussion

**Observation:** The speedup improves up to 8 threads ( $S_8 \approx 2.47$ ) but fails to scale linearly. Notably, at 16 threads, the performance actually **degrades** (25.69s vs 23.67s at 8 threads).

#### Root Causes:

- **Memory Bound Algorithm:** The Jacobi method involves dense matrix-vector multiplication. The CPU is likely waiting for data from RAM rather than performing calculations. Adding more threads just increases contention for memory bandwidth, saturating the bus.
- **Parallel Overhead:** At 16 threads, the cost of managing threads (creation, synchronization barriers at the end of loops) outweighs the benefit of splitting the work further.
- **False Sharing / Cache Issues:** Although ‘x\_courant’ is accessed distinctly by index  $i$ , the frequent reads of the shared matrix  $a$  and vector  $x$  across all cores cause cache thrashing.